

AI Interview Evaluation Report

Candidate	Sunil
Email	abhay0005@gmail.com
Target Role	senior AI
Experience	13
Section	Generative AI / LLM fundamentals
LLM	Ollama (Local) / llama3.2:1b
Status	completed
Overall Score	2.43/10
Overall Feedback	The candidate showed a foundation that needs further preparation (overall score: 2.4/10). Focus on deepening understanding of LLM fundamentals and practical applications.
Questions Attempted	7

Question 1: What is an LLM, and how does it work?

Candidate Answer: A Large Language Model (LLM) is a type of artificial intelligence program designed to read, write, and understand human language. It works as an advanced statistical prediction machine, processing text by breaking it down into small pieces and repeatedly guessing the most likely next word in a sequence using deep learning and neural networks.

Score	4.0	Correctness		Completeness	
Depth	6.0	Clarity		Response Time	27.073161363601685

Feedback:

Question 2: What is a Transformer architecture?

Candidate Answer: The transformer architecture is a deep learning design introduced in the 2017 paper "Attention Is All You Need". It processes entire sequences of data in parallel using self-attention mechanisms rather than step-by-step loops, powering modern AI systems like ChatGPT, Gemini, and Llama

Score	2.0	Correctness		Completeness	
Depth	5.0	Clarity		Response Time	19.746129751205444

Feedback:

Question 3: Explain attention and self-attention.

Candidate Answer: Attention and self-attention are core tools in modern artificial intelligence, especially in Large Language Models and Transformers. Attention lets a model look at different parts of an input sequence to find relevant information. Self-attention is a specific type of attention where a single sequence looks at its own tokens to understand context and relationships between words

Score	2.0	Correctness		Completeness	
Depth	7.0	Clarity		Response Time	21.9217529296875

Feedback:

Question 4: What are tokens, embeddings and context windows?

Candidate Answer: Tokens are the basic building blocks of text that AI models process, such as words or parts of words. Embeddings convert those text tokens into lists of numbers that capture meaning and relationships. A context window is the maximum number of input and output tokens an AI can remember during a single

conversation or task.

Score	2.0	Correctness		Completeness	
Depth	5.0	Clarity		Response Time	22.09917688369751

Feedback:

Question 5: What is the difference between GPT, Llama and Claude?

Candidate Answer: GPT (by OpenAI), Claude (by Anthropic), and Llama (by Meta) are leading large language models with core differences in ownership and accessibility. GPT and Claude are closed-source, commercial services accessed via APIs or chat apps, whereas Llama is an open-weight model that users can download, run, and modify locally.

Score	2.0	Correctness		Completeness	
Depth	6.0	Clarity		Response Time	26.206339836120605

Feedback:

Question 6: What is temperature? What happens when you increase it?

Candidate Answer: Temperature is a measure of the average speed and energy of the tiny particles (atoms and molecules) that make up an object or substance. When you increase the temperature, these particles gain energy, move or vibrate faster, and push slightly farther apart.

Score	3.0	Correctness		Completeness	
Depth	9.0	Clarity		Response Time	23.367022037506104

Feedback:

Question 7: What are hallucinations and how can you reduce them?

Candidate Answer: Hallucinations are false or fabricated responses presented as absolute facts. In humans, they involve sensing things that do not exist. In artificial intelligence, large language models generate plausible-sounding untruths due to statistical word prediction gaps

Score	2.0	Correctness		Completeness	
Depth	5.0	Clarity		Response Time	29.851197004318237

Feedback: